Model choice moves the number far less than the features do

Every configuration we ran, on CB-full (231 apo structures). One row per condition — same features, same negatives — so the horizontal spread inside a row is the model effect alone (0.003–0.030 AUPRC). The vertical distance between rows is the feature effect (0.398 → 0.476). Axes are cropped to the model range: a dot's position encodes the value, and the full 0–1 range would collapse every dot in a row onto one pixel.

Our sklearn controls exist on CB-full only — the stored artifacts keep summary metrics, not per-residue scores, so they cannot be re-scored on CB-P2RANK-apo. · TabICL has no full-304k row by construction: 304k rows do not fit in context, so the full-304k rows are the sklearn models' best case — and TabICL in 50k of context still leads them. · RSA alone, the single-feature floor, sits far below the cropped axes at AUPRC 0.089 / AUC 0.499 / MCC 0.034. · Rebuild: python scripts/model_ladder.py